

# Message Passing Interface

## MPhil in Scientific Computing, 2022-2023

Maria Nikodemou

Centre for Scientific Computing  
University of Cambridge

# Outline

- 1 Parallel programming for scientific computing
- 2 Message Passing Interface
- 3 MPI: The basics
- 4 Point-to-Point Communication
- 5 Collective communication
- 6 More Point-To-Point
- 7 Error handling
- 8 Common mistakes
- 9 Performance Tips
- 10 Further topics

## Books

- W. Gropp et al. *Using MPI: portable parallel programming with the message-passing interface*, Vol. 1 (MIT Press, 1999) - thorough coverage of MPI
- Peter Pacheco, *An introduction to parallel programming* (Elsevier, 2011) - introductory
- R. Robey and Y. Zamora, *Parallel and High Performance Computing* (Simon and Schuster, 2021) - introductory

## MPI Standard

- <http://www.mpi-forum.org/docs/> - MPI Standard

# About this course

## Prerequisites

- Knowledge of C/C++ or Fortran
- Access to a multi-core computer

## This course

- aims to show you how to safely implement MPI in your code
- covers a small part of MPI, sufficient for most applications
- provides MPI functions and variables in both C and Fortran, but most examples are given in C++ only
- is based upon material provided by Dr. Rutter, who previously lectured this course

# Parallel programming for scientific computing

# What is parallel computing

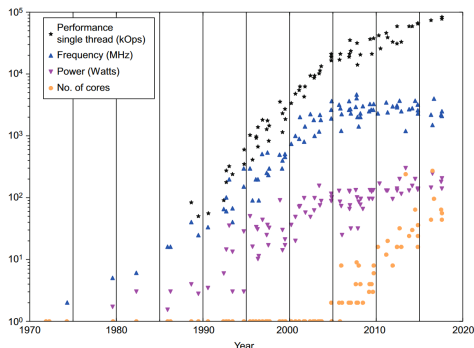
## Parallel computing

The simultaneous execution of multiple operations

Can be achieved through

- process-based parallelisation (e.g. MPI)
- thread-based parallelisation (e.g. OpenMP)
- stream processing (e.g. CUDA)
- vectorisation

# Why parallel?



Microprocessor trend data (Robey and Zamora, *Parallel and High Performance Computing*)

Before 2005,

- only single core CPUs existed
- clock frequency (instruction execution speed) kept increasing

From 2005 onward,

- clock frequency flattens out
- number of cores keep increasing

Optimal CPU performance is only achievable through parallel computing!


# Benefits of parallel computing

- **Speedup**: reduction of run time by orders of magnitude
- **Bigger scales**: ability to reach problem dimensions not possible with serial codes
- **Energy efficiency**: reduction of power consumption (and costs) and extension of battery life




# Distributed vs. Shared Memory Architecture

## Distributed memory architecture



- each processor has its own local memory
- processors exchange messages with each other through a communication network

## Shared memory architecture



- processors share the same address space
- processors have access to the pool of shared memory

# Terminology

## Thread

A separate instruction pathway through a process

## Process

An independent unit of computation that has ownership of a portion of memory and control over resources in user space

## Rank

A unique identifier (usually an integer) to distinguish the individual process within the set of processes

## Core

The basic element of the system that does the mathematical and logical operations

## Computer processing unit (CPU), or processor

The discrete processing device composed of one or more computational cores that is placed on the socket of a circuit board to provide the main computational operations


## Node

A basic building block of a compute cluster with its own memory (accessible by all its processors) and a network to communicate with other compute nodes

Definitions from Robey and Zamora, *Parallel and High Performance Computing*


# Distributed vs. Shared Memory Architecture

## Distributed memory architecture



- + Scalable: limitless ability to incorporate more nodes
- Needs more thought:
  - data distribution among the different memory regions
  - keeping process communication to a minimum

## Shared memory architecture



- + Simplified programming
- Need for memory access synchronisation to avoid potential memory conflicts
- Limited scalability:
  - adding more CPUs does not increase the amount of memory available
  - large core-count machines are rare and expensive
- Required process communication is hidden at the hardware level (so less obvious to tune)

# Process vs. Thread based parallelisation


## Process-based parallelisation

- Developed for distributed memory architectures
- The program creates separate processes, or *ranks*, with their own memory space
- Messages are used to move data between processes explicitly
- The OS distributes the processes on the processors and can move these during run time
- The **Message Passing Interface** (MPI) dominates process-based parallel programming

## Thread-based parallelisation

- Developed for shared memory architectures
- The program creates separate instruction paths, or *threads*, within the same process
- The process memory can be easily shared between threads
- The OS decides where to place the threads
- The **OpenMP** standard dominates thread-based parallel programming

# Process vs. Thread based parallelisation



Conceptual visualization of the programming effort required to improve performance using either MPI or OpenMP (Robey and Zamora, *Parallel and High Performance Computing*)

# Message Passing Interface

# What is MPI?

## Message Passing Interface (MPI)

A **message-passing library interface specification**, designed for running jobs of multiple processes on multi-core computers and on multiple computers.

- MPI is a library, not a language
- It addresses the message-passing parallel programming model
- It provides the appropriate functions, subroutines or methods to move data from the address space of one process to that of another process
- As such, MPI can be used to parallelise existing, or new code

# History of MPI

- May 1994 • **Version 1.0:** the first MPI standard, with participation from over 40 vendor-neutral organisations internationally. The goal was to develop a widely used standard for writing message-passing programs
- Jul 1997 • **Version 2.0:** corrections and extensions to MPI-1.0
- Sep 2012 • **Version 3.0:** significant extensions to MPI functionality
- Jun 2021 • **Version 4.0:** major update with addition of new features



MPI is **not** a language. MPI operations are expressed as functions, subroutines, or methods, according to the appropriate language bindings:

- MPI was originally developed for Fortran and C languages
- Fortran:
  - originally, MPI-1.1 provided bindings for F77
  - MPI can still be used with most Fortran 77 compilers
  - now, MPI bindings are for Fortran 90 or later
- C++ programmers can use MPI through C bindings
  - C++ bindings were introduced in MPI-2.0 but deprecated in MPI-2.2 and removed in MPI-3.0 - these are not recommended

# MPI implementations

- MPI is a specification, not an implementation
- Various implementations of the MPI standard exist
  - offered by parallel computer vendors for their machines, or
  - free, publicly available on the Internet
- The two major open source implementations of MPI are
  - **OpenMPI**: [www.open-mpi.org](http://www.open-mpi.org)
  - **MPICH**: [www.mpich.org](http://www.mpich.org)
- OpenMPI is the more widely used free implementation
- Many commercial MPIs are based on MPICH
- These will conform to the MPI standard, but may differ in matters not specified by the standard, like
  - the initial loading of the executables onto the parallel machine
  - which processes are allowed to execute input/output tasks
- It is very important to **use the same MPI implementation** at compile time, link time and run time!

## MPI: The basics

# Compiling an MPI program

- MPI is a library, not a language
- No special compiler or accommodations for the OS are required
- You can compile code as usual, with the use of flags to link in the MPI library
- Or use a **compiler wrapper** which ensures that all required libraries and options are properly applied:

Language	Wrapper compiler name
C	mpicc
C++	mpicxx, mpiCC, mpic++
Fortran	mpifort, mpif90, mpif77

- For example,  
`$ mpicxx my_program.cpp -o my_program`
- Other/more compiler wrappers are defined by different MPI vendors, so one should check the corresponding MPI library manual (e.g. `mpiicc`, `mpiicpc` and `mpiifort` for Intel MPI)

# Running an MPI program

- To run an MPI program, use  
`$ mpiexec -n <number of processes> ./my_program`
- The startup command for MPI programs is not part of the MPI standard - it may vary between implementations
- The MPI Forum recommends the use of `mpiexec`
- Most implementations also support the syntax  
`$ mpirun -np <number of processes> ./my_program`
- You should generally use no more processes than the number of CPU cores in your machine.
- Note: if you use more processes than CPU cores, `mpiexec` may produce an error, depending on the MPI implementation. For example, in OpenMPI you can use the `-oversubscribe` option to get around this.

# Code structure - Hello parallel world

## helloWorld.cpp

```
#include <iostream>
#include <mpi.h>
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    std::cout << "Hello parallel world" << std::endl;

    MPI_Finalize();
    return 0;
}
```

## helloWorld.f08

```
program hello
    use mpi_f08

    call mpi_init()

    write(*,*) 'Hello parallel world'

    call mpi_finalize()

end program hello
```

- These are simply C++/Fortran codes, with the addition of a few lines of MPI-related code:
  - The header file `mpi.h` or module `mpi_f08` provides definitions of named constants and types, and function prototypes. Fortran: You may also see `use mpi` or `include 'mpif.h'` in place of `use mpi_f08`, but these are now discouraged by the standard
  - The call to `MPI_Init` for MPI initialisation
  - The call to `MPI_Finalize` for MPI completion

# Hello parallel world

```
$ mpicxx helloWorld.cpp
```

```
$ mpiexec -n 4 ./a.out
```

```
Hello parallel world
```

```
Hello parallel world
```

```
Hello parallel world
```

```
Hello parallel world
```

```
$ mpifort helloWorld.f08
```

```
$ mpiexec -n 3 ./a.out
```

```
Hello parallel world
```

```
Hello parallel world
```

```
Hello parallel world
```

- All identifiers defined by MPI start with the string MPI\_.
- C/C++:
  - MPI function names and MPI-defined types: The first letter after MPI\_ is capitalised
  - MPI-defined macros and constants: All letters are capitalised
- To avoid confusion, do not use the prefix MPI\_ (or even PMPI\_) in your variable and function names!
- In general, MPI functions in C/C++ return an int, and Fortran MPI routines have an ierror argument, which contains the error code. When the call is successful, this gives MPI\_SUCCESS (which is equal to 0)
  - Fortran: with the mpi\_f08 module, this argument is declared as optional
- The most basic MPI functions are MPI\_Init and MPI\_Finalize
- These should be called exactly once in any MPI program



# MPI\_Init(): Initialising MPI

```
int MPI_Init(int* argc_p, char*** argv_p);
```

```
MPI_Init(ierror)
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- This sets up the MPI environment, e.g. decides which process gets which rank, allocates storage for message buffers
- The call to MPI\_Init should precede all other MPI function calls <sup>1</sup>
- This **does not** mark the point at which the program moves from begin serial to parallel!
- mpiexec launches the specified number of processes from the beginning, and this number is fixed for the duration of the MPI program's execution
- C/C++: MPI\_Init accepts pointers to the argc and argv arguments of the main function, but also accepts NULL

---

<sup>1</sup>With a few exceptions, like MPI\_Get\_version and MPI\_Initialized

# MPI\_Finalize(): Finalising MPI

```
int MPI_Finalize(void);
```

```
MPI_Finalize(ierr)
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierr
```

- This terminates the MPI environment: it cleans up all MPI data-structures
- The call to MPI\_Finalize should follow all other MPI function calls <sup>2</sup>
- Unless the program is aborted, all MPI processes should call MPI\_Finalize
- MPI\_Finalize must be called when all communications have completed - this is the programmer's responsibility

---

<sup>2</sup>With a few exceptions, like MPI\_Get\_version and MPI\_Finalized

## MPI\_Get\_version(): MPI standard version

```
int MPI_Get_version(int *version, int *subversion);
```

```
MPI_Get_version(version, subversion, ierror)  
    INTEGER, INTENT(OUT) :: version, subversion  
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- This is used to determine which version of the MPI Standard is in use
- The output is stored in the variables version and subversion
- So, for example version = 4 and subversion = 0 corresponds to MPI 4.0
- This is one of the few functions that can be called at any time in an MPI program - even before MPI\_Init or after MPI\_Finalize

# Hello parallel world: size and rank

## helloWorld\_ranks.cpp

```
#include <iostream>
#include <mpi.h>
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    std::cout << "Hello parallel world, I am process "\
        << rank << " out of " << size << "\n";

    MPI_Finalize();
    return 0;
}
```

## helloWorld\_ranks.f08

```
program hello

    use mpi_f08

    integer :: size, rank

    call mpi_init()
    call mpi_comm_size(mpi_comm_world, size)
    call mpi_comm_rank(mpi_comm_world, rank)

    write(*, 101) rank, size
101 format('Hello parallel world, I am process ', &
    I2, ' out of ', I2)

    call mpi_finalize()

end
```

# Hello parallel world: size and rank

```
$ mpicxx helloWorld_ranks.cpp
$ mpiexec -n 8 ./a.out
Hello parallel world, I am process 2 out of 8
Hello parallel world, I am process 3 out of 8
Hello parallel world, I am process 4 out of 8
Hello parallel world, I am process 6 out of 8
Hello parallel world, I am process 7 out of 8
Hello parallel world, I am process 0 out of 8
Hello parallel world, I am process 1 out of 8
Hello parallel world, I am process 5 out of 8
```

- **Non-deterministic output:** processes compete for access to the shared output device, resulting in unpredictable output order
- This order might change each time you run the program
- We could make all other processes send their message to process 0, and process 0 print the output in process rank order

- MPI has the concept of **communicators**, to provide the appropriate scope for all communication operations
- Communicators can be thought as collections of processes that can send messages to each other
- All MPI communication function calls require a communicator argument
- MPI\_COMM\_WORLD is the **default communicator**, which consists of all the processes in the program - this is defined in `mpi.h` and `mpi_f08`
- It is possible to create user-defined communicators - this is more advanced, but useful for multiple levels of parallelism

# Communicator size and rank

- **Communicator size:** the total number of processes in the communicator. This is stored in the variable `size` after calling:

```
int MPI_Comm_size(MPI_Comm comm, int *size);
```


```
MPI_Comm_size(comm, size, ierror)  
  TYPE(MPI_Comm), INTENT(IN) :: comm  
  INTEGER, INTENT(OUT) :: size  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- **Process rank:** the unique identifier assigned to the individual process within the communicator. This is an integer between 0 and `size-1`, obtained in variable `rank` after calling:

```
int MPI_Comm_rank(MPI_Comm comm, int *rank);
```

```
MPI_Comm_rank(comm, rank, ierror)  
  TYPE(MPI_Comm), INTENT(IN) :: comm  
  INTEGER, INTENT(OUT) :: rank  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

# Communicator size and rank





# MPI\_Abort(): Aborting MPI programs

```
int MPI_Abort(MPI_Comm comm, int errorcode);
```

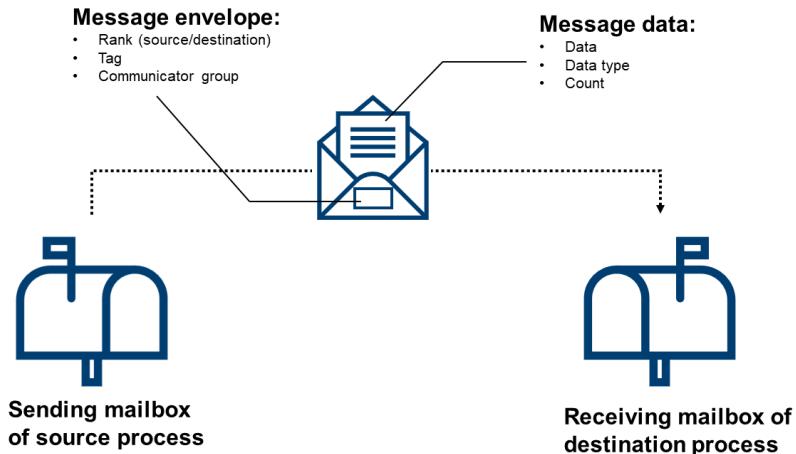
```
MPI_Abort(comm, errorcode, ierror)  
  TYPE(MPI_Comm), INTENT(IN) :: comm  
  INTEGER, INTENT(IN) :: errorcode  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- To abort all processes in an MPI program, run this with MPI\_COMM\_WORLD as the comm argument
- the errorcode is returned from the main program to the invoking environment
- For example, the following will abort because of too few processes:

```
if ((rank == 0) && (nproc < 2)){  
    std::cerr << "Too few processes\n" << std::endl;  
    MPI_Abort(MPI_COMM_WORLD, 1);  
}
```

## Point-to-Point Communication

# Send/Receive message structure



# Send/Receive message structure

- Passing a message from one process to another involves two operations:
  - **Sending** the item from the source process
  - **Receiving** the item from the destination process
- At each operation, the following is specified
  - the **pointer to memory** where the message can be found, or will be stored
  - the message **size**
  - the message **data type**
  - the **rank** of the source/destination process
  - the **tag** (a non-negative integer used to distinguish messages)
  - the **communicator** within which the message is sent

# MPI\_Send(): Send operation

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest,\n             int tag, MPI_Comm comm);
```

```
MPI_Send(buf, count, datatype, dest, tag, comm, ierror)\n  TYPE(*), DIMENSION(..), INTENT(IN) :: buf\n  INTEGER, INTENT(IN) :: count, dest, tag\n  TYPE(MPI_Datatype), INTENT(IN) :: datatype\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- buf: initial address of send buffer
- count: number of elements in send buffer
- datatype: datatype of each send buffer element
- dest: rank of destination
- tag: message tag
- comm: communicator

## MPI\_Recv(): Receive operation

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source,\n             int tag, MPI_Comm comm, MPI_Status *status)
```

```
MPI_Recv(buf, count, datatype, source, tag, comm, status, ierror)\nTYPE(*), DIMENSION(..) :: buf\nINTEGER, INTENT(IN) :: count, source, tag\nTYPE(MPI_Datatype), INTENT(IN) :: datatype\nTYPE(MPI_Comm), INTENT(IN) :: comm\nTYPE(MPI_Status) :: status\nINTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- buf: initial address of receive buffer
- count: number of elements in receive buffer
- datatype: datatype of each receive buffer element
- dest: rank of source
- tag: message tag
- comm: communicator
- status: status object

# Example: One-way communication

## oneWay\_comm.cpp

```
#include <iostream>
#include <mpi.h>
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    MPI_Status status;

    double x;
    if (rank == 0)
    {
        x = 10.0;
        MPI_Send(&x, 1, MPI_DOUBLE, 1, 0, MPI_COMM_WORLD);
    }
    else if (rank == 1)
    {
        MPI_Recv(&x, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &
status);
        std::cout << "Rank 1 has received x = " << x << \
" from rank 0" << std::endl;
    }

    MPI_Finalize();
    return 0;
}
```

## oneWay\_comm.f08

```
program hello

    use mpi_f08

    integer :: rank
    real :: x
    type(MPI_Status) :: status

    call mpi_init()
    call mpi_comm_rank(mpi_comm_world, rank)

    if (rank .eq. 0) then
        x = 10.0
        call MPI_Send(x, 1, MPI_REAL, 1, 0, MPI_COMM_WORLD)
    else if (rank .eq. 1) then
        call MPI_Recv(x, 1, MPI_REAL, 0, 0, MPI_COMM_WORLD, &
status)
        write(*, 101) x
101 format('Rank 1 has received x = ', F4.1, &
' from rank 0')
    end if

    call mpi_finalize()

end
```

- For the message to be successfully received, we need
  - the same communicator in both calls
  - the sending/receiving processes to match those specified in the receive/send calls, respectively
  - the same tag in both calls
  - the same data types in both calls
  - the receiving buffer size to be at least as big as the message size
- The tag should be between 0 and 32767, but need not be unique for each message



- The wildcard `MPI_ANY_TAG` can be used in the receive call if an explicit number is not desired
- The wildcard `MPI_ANY_SOURCE` can be specified in the receive call to match data coming from any source within the specified communicator group
- No wildcards are available for `MPI_Send` - the sender must specify both a destination rank and a tag

- A process can receive a message without knowing the message's sender, tag or size
- The above details are returned in the status argument of `MPI_Recv` and obtained using:
  - source process: `status.MPI_SOURCE` or `status%MPI_SOURCE`
  - message tag: `status.MPI_TAG` or `status%MPI_TAG`
  - the number of entries received (count of datatype, not bytes!):  
`MPI_Get_count(status, datatype, count)`
- If these details are not required, one may pass the predefined constant `MPI_STATUS_IGNORE` as an argument of the receive function

# MPI data types

- It is important to always specify the correct datatypes for messages
- MPI provides a large set of predefined datatypes, including all basic datatypes in C and Fortran


MPI datatype	C datatype	MPI datatype	Fortran datatype
MPI_INT	signed int	MPI_INTEGER	INTEGER
MPI_FLOAT	float	MPI_REAL	REAL
MPI_DOUBLE	double	MPI_DOUBLE_PRECISION	DOUBLE PRECISION
MPI_CHAR	char	MPI_COMPLEX	COMPLEX
MPI_C_BOOL	_Bool	MPI_CHARACTER	CHARACTER(1)
MPI_LONG	signed long int	MPI_LOGICAL	LOGICAL

- If using C++, the MPI datatype for a bool is MPI\_CXX\_BOOL

# Blocking Communication

- The MPI\_Send and MPI\_Recv are **blocking** calls
- MPI\_Send only returns when the message data and envelope are safely stored away and the contents of the send buffer can be safely modified without affecting the data being sent
- The message might be copied directly into the matching receive buffer, or temporarily stored into a system buffer
- It is up to MPI to decide whether sent messages will be buffered or not (based on available buffer space and performance)
- Therefore, MPI\_Send may return before or after the call to a matching receive
- MPI\_Recv only returns when the message is received

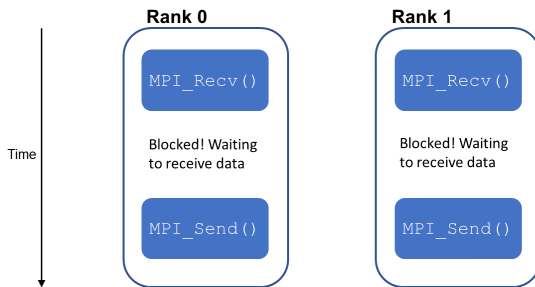
# Message exchange between two processes



- Consider a simple program with two processes sending a message to each other
- Which operation should be called first?

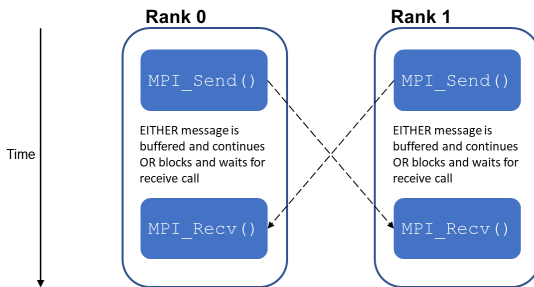
# Receive first

- Both processes call `MPI_Recv` first
- In both processes, the send would not be called until after the receive completes
- But the receive cannot complete if a send is not called from the other process
- This causes the program to hang, also known as **deadlock**



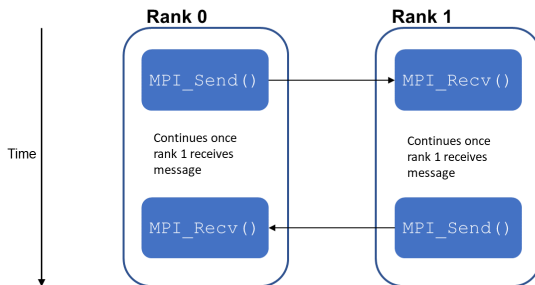
# Send first

- Both processes call `MPI_Send` first
- MPI will either copy the message to a temporary buffer or wait for the receive call
- If none of the two messages is buffered (e.g. if the messages are large), the send calls will not complete and program will hang
- Therefore, calling the send function first is not safe either



# Alternate send/receive

- Rank 0 calls MPI\_Send first and Rank 1 calls MPI\_Recv first
- The first message arrives, allowing operation of the second send-receive pair
- The messages are safely passed and deadlock is avoided!





## Example: alternate send/receive

- Alternate send/receive can be implemented as follows:

```
if (rank == 0)
{
    MPI_Send(sendbuf, sendcount, sendtype, 1, sendtag, comm);
    MPI_Recv(recvbuf, recvcount, recvtype, 1, recvtag, comm, status);
}
else if (rank == 1)
{
    MPI_Recv(recvbuf, recvcount, recvtype, 0, recvtag, comm, status);
    MPI_Send(sendbuf, sendcount, sendtype, 0, sendtag, comm);
}
```

- This can get complicated when dealing with more processes and more complex communication
- One way round this: the MPI\_Sendrecv function

# MPI\_Sendrecv(): Send and receive operations

```
int MPI_Sendrecv(const void *sendbuf, int sendcount, MPI_Datatype sendtype,\n                 int dest, int sendtag, void *recvbuf, int recvcount, MPI_Datatype recvtype,\n                 int source, int recvtag, MPI_Comm comm, MPI_Status *status);
```

```
MPI_Sendrecv(sendbuf, sendcount, sendtype, dest, sendtag, recvbuf, recvcount, &\n             recvtype, source, recvtag, comm, status, ierror)
```

```
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf
```

```
INTEGER, INTENT(IN) :: sendcount, dest, sendtag, recvcount, source, recvtag
```

```
TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype
```

```
TYPE(*), DIMENSION(..) :: recvbuf
```

```
TYPE(MPI_Comm), INTENT(IN) :: comm
```


```
TYPE(MPI_Status) :: status
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Executes a send-receive operation: sends a message to one destination and receives another message from another (or the same) process
- Both send and receive use the same communicator, but possibly different tags, counts and datatypes
- The send and receive buffers should be different
- MPI\_Sendrecv effectively places the responsibility for correct communication execution to the MPI library, avoiding deadlocks

# Send-Receive: circular chain


- We've seen how a call to `MPI_Sendrecv` can be used for communication between two processes
- This can be easily extended to larger circular chains, for example



```
MPI_Sendrecv(sendbuf, sendcount, sendtype, (rank + 1) % size, sendtag, recvbuf, \
recvcount, recvttype, (rank + size - 1) % size, recvttag, comm, status);
```

# Send-Receive: multiple chains

```
MPI_Sendrecv(sendbuf, sendcount, sendtype, dest, sendtag, recvbuf, recvcount, \
recvtype, source, recvtag, comm, status);
```



Rank	dest	source
0	1	MPI_PROC_NULL
1	2	0
2	MPI_PROC_NULL	1
3	4	4
4	3	3
5	MPI_PROC_NULL	MPI_PROC_NULL

- MPI\_PROC\_NULL is a dummy process
- When used as a destination or source argument in a send/receive call, then the communication has no effect - i.e. the call succeeds and returns as soon as possible without taking any action
- Example use: to simplify code for dealing with boundaries

# MPI\_Sendrecv\_replace(): Send and receive with message replacement

```
int MPI_Sendrecv_replace(void *buf, int count, MPI_Datatype datatype, int dest,\n                          int sendtag, int source, int recvtag, MPI_Comm comm, MPI_Status *status);
```

```
MPI_Sendrecv_replace(buf, count, datatype, dest, sendtag, source, recvtag,&\n                     comm, status, ierror)\n  TYPE(*), DIMENSION(..) :: buf\n  INTEGER, INTENT(IN) :: count, dest, sendtag, source, recvtag\n  TYPE(MPI_Datatype), INTENT(IN) :: datatype\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  TYPE(MPI_Status) :: status\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Executes a blocking send-receive operation, with the same buffer used for both the send and the receive
- So the message sent is replaced by the message received
- For example, if used in a circular chain, this will simply permute messages around

## Example application: Estimating $\pi$

- Let's apply what we learnt so far to something more useful: calculating the value of  $\pi$ , using the Leibniz formula:

$$\pi = 4 \sum_{i=0}^{\infty} \left( (-1)^i \frac{1}{2i+1} \right) = 4 \left[ 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right]$$

- We will truncate the infinite sum to some integer  $n$ :

$$\pi \approx 4 \sum_{i=0}^n \left( (-1)^i \frac{1}{2i+1} \right) = S$$

# Summation code: serial

## serial\_sum.cpp

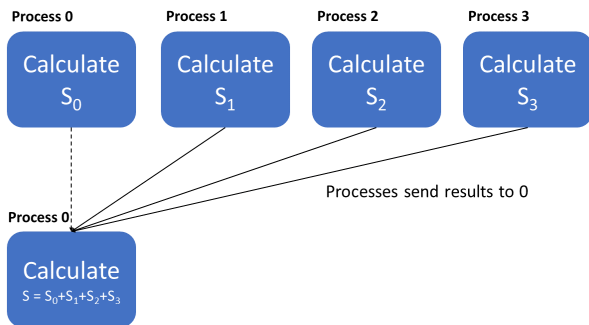
```
#include <iostream>
#include <iomanip>
#include <cmath>

int main(int argc, char *argv[])
{
    long n = std::stol(argv[1]);
    double sum = 0.0;
    for (long i = 0; i < n; i++)
    {
        double recip = 1.0/(2.0 * i + 1.0);
        if (i%2 == 0)
        {
            sum += recip;
        }
        else
        {
            sum -= recip;
        }
    }
    sum *= 4.0;

    std::cout << std::setprecision(10);
    std::cout << "Pi ~ " << sum << " with error " << M_PI - sum << std::endl;
    return 0;
}
```

# Parallelisation

- Increasing  $n$  gives greater accuracy in our approximation
- Each term in the sum is independent of the others, so we can parallelise
- A possible way to parallelise with  $n_{\text{proc}}$  processes:
  - Split the total sum,  $S$ , into  $n_{\text{proc}}$  sums:  $S = S_0 + S_1 + \dots + S_{n_{\text{proc}}-1}$
  - Process  $i$  calculates  $S_i$  and sends it to process 0
  - Process 0 computes  $S$  by adding the  $S_i$  together





# Summation code: parallel

## parallel\_sum.cpp

```
#include <iostream>
#include <iomanip>
#include <cmath>
#include <mpi.h>

int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    long n = std::stol(argv[1]);

    int rank, size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    double sum = 0.0;
    long i_start = n * rank / size;
    long i_end = n * (rank + 1) / size;

    for (long i = i_start; i < i_end; i++)
    {
        double recip = 1.0/(2.0 * i + 1.0);
        if (i%2 == 0)
        {
            sum += recip;
        }
        else
        {
            sum -= recip;
        }
    }
}
```

```
//Continued code
if (rank != 0)
{
    //Send result to rank 0 (tag=0)
    MPI_Send(&sum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}
else
{
    double global_sum = sum;
    for (int i = 1; i < size; i++)
    {
        double local_sum;
        //Receive results from all other ranks (tag=0)
        MPI_Recv(&local_sum, 1, MPI_DOUBLE, \
        MPI_ANY_SOURCE, 0, MPI_COMM_WORLD, \
        MPI_STATUS_IGNORE);

        //Update global sum
        global_sum += local_sum;
    }
    global_sum *= 4.0;

    std::cout << std::setprecision(10);
    std::cout << "Pi ~ " << global_sum << " with error "\
    << M_PI - global_sum << std::endl;
}

MPI_Finalize();

return 0;
}
```

# Results

```
$ g++ serial_sum.cpp -o serial_sum
$ ./serial_sum 10
Pi ~ 3.041839619 with error 0.09975303466
$ ./serial_sum 10000
Pi ~ 3.141492654 with error 9.999999976e-05
$ ./serial_sum 10000000
Pi ~ 3.141592554 with error 1.000000016e-07

$ mpicxx parallel_sum.cpp -o parallel_sum
$ mpiexec -n 4 parallel_sum 10
Pi ~ 3.041839619 with error 0.09975303466
$ mpiexec -n 4 parallel_sum 10000
Pi ~ 3.141492654 with error 9.999999975e-05
$ mpiexec -n 4 parallel_sum 10000000
Pi ~ 3.141592554 with error 1.000000505e-07
```

## Note

We get slightly different errors when running on more processes. This is expected. It is because the order of floating point summation changes when introducing more processes

# Summation code performance

- The simplest way to time our program is to use the Unix shell command `time`
- This outputs the time taken to run a program from start to finish.



```
$ time ./serial_sum 10000000
Pi ~ 3.141592554 with error 1.000000016e-07
real    0m0.028s
user    0m0.026s
sys     0m0.000s
```

tells us that the program runs in 28ms

- Timing a command only once may give misleading results. One could use `multitime` instead:

```
$ multitime -n 2 ./serial_sum 10000000
Pi ~ 3.141592554 with error 1.000000016e-07
Pi ~ 3.141592554 with error 1.000000016e-07
===> multitime results
1: ./serial_sum 10000000
Mean      Std.Dev.   Min         Median      Max
real      0.029      0.002      0.027      0.029      0.031
user      0.028      0.002      0.027      0.028      0.030
sys       0.000      0.000      0.000      0.000      0.000
```



# Serial vs. parallel



- It appears to take longer to run in parallel (with 4 processes) than serially! Why?
- Because of the **parallel overhead**, typically coming from communication

$$T_{parallel}(n, p) = T_{serial}(n)/p + T_{overhead}(n, p)$$


# Serial vs. parallel




- The efforts of parallelisation only pay off when the problem size is big enough
- The parallel overhead does increase with the problem size, but at a much slower rate compared to the execution time
- For  $n = 10^{11}$ , we get a speed up  $S(n, p) = \frac{T_{serial}(n)}{T_{parallel}(n, p)}$  of  $\sim 3.8$

# Number of processes

$$n = 10^{11}$$




$$n = 10^4$$




- For a large problem size, run time decreases with more processes, as one would expect
- This is not the case for a smaller problem size. When using more processes, the communication overheads become increasingly significant and cause an overall increase in run time

# Simple structure



- In our algorithm to estimate  $\pi$ , each process calculated a value which was then sent to process 0 to compute the global sum
- While process 0 is calculating this sum, the other processes are doing nothing
- We could try to distribute the workload among the processes more evenly

# Tree-structure



- Iterations of odd-positioned processes sending their results to even-positioned processes to do the sum
- Assuming that the message-passing and calculations occur simultaneously at each stage, the scheme requires 2 receives and 2 additions (compared with 3 receives and 3 additions in previous scheme)
- If using 16 processes, this improves the original scheme by a factor of 4



# Which communication structure?

- A tree-like communication structure is much harder to implement than the original, simple structure
- There are many other possible candidates
- The best structure might depend on the number of processes we use, or even worse, the problem we're solving
- How do we choose the optimal one?
- We don't have to: MPI's collective communication functions come to the rescue

## Collective communication

# Collective communication

- Collective communication refers to communication that involves a group or groups of processes
- MPI provides a number of collective functions
- Collective functions involve **all** processes within a communicator
- They can do two kinds of operations:
  - Data distribution, e.g. broadcasting a message to all processes
  - Collective computation, e.g. calculating the minimum, maximum, sum etc.
- These effectively remove the responsibility of optimising communication from the programmer and place it on the MPI implementation developer
- They should be used in place of point-to-point communication, whenever possible

# MPI\_Barrier(): Process synchronisation


```
int MPI_Barrier(MPI_Comm comm);
```

```
MPI_Barrier(comm, ierror)
```


```
TYPE(MPI_Comm), INTENT(IN) :: comm
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- **Barrier** is the simplest collective operation: it synchronises the processes within the communicator, comm
- No process can return from the call to MPI\_Barrier before all processes have entered
- This is the only collective operation which guarantees synchronisation!
- Note: this does not guarantee that all processes will exit the call at the same time



# Reduction operation



- **Reduction:** one of the most commonly used collective function
- It takes values from all processes within the communicator and combines them into a scalar result, according to the operator:

MPI reduction operator	Meaning
MPI_MAX / MPI_MIN	Maximum / Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND / MPI_LOR	Logical AND / OR
MPI_BAND / MPI_BOR	Bitwise AND / OR
MPI_LXOR	Logical exclusive OR
MPI_BXOR	Bitwise exclusive OR
MPI_MAXLOC / MPI_MINLOC	Maximum / minimum and location

- It is also possible to define your own reduction operator

## MPI\_Reduce(): Reduction operation

```
int MPI_Reduce(const void *sendbuf, void *recvbuf, int count, MPI_Datatype\
               datatype, MPI_Op op, int root, MPI_Comm comm);
```

```
MPI_Reduce(sendbuf, recvbuf, count, datatype, op, root, comm, ierror)
    TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf
    TYPE(*), DIMENSION(..) :: recvbuf
    INTEGER, INTENT(IN) :: count, root
    TYPE(MPI_Datatype), INTENT(IN) :: datatype
    TYPE(MPI_Op), INTENT(IN) :: op
    TYPE(MPI_Comm), INTENT(IN) :: comm
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- The send and receive buffers must be different to avoid **aliasing**
- The receive buffer is only relevant in the root (destination) process - it is a dummy buffer in all other processes
- By setting  $\text{count} > 1$ , MPI\_Reduce effectively operates on arrays, instead of scalars

# Estimating $\pi$ - revisited

## parallel\_reduce\_sum.cpp

```
#include <iostream>
#include <iomanip>
#include <cmath>
#include <mpi.h>

int main(int argc, char *argv[])
{
    MPI_Init(NULL, NULL);

    long n = std::stol(argv[1]);

    int rank, size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    double sum = 0.0;
    long i_start = n * rank / size;
    long i_end = n * (rank + 1) / size;

    for (long i = i_start; i < i_end; i++)
    {
        double recip = 1.0/(2.0 * i + 1.0);
        if (i%2 == 0)
        {
            sum += recip;
        }
        else
        {
            sum -= recip;
        }
    }
}
```

```
//Continued code
double global_sum;
MPI_Reduce(&sum, &global_sum, 1, MPI_DOUBLE, \
    MPI_SUM, 0, MPI_COMM_WORLD);

if (rank == 0)
{
    global_sum *= 4.0;
    std::cout << std::setprecision(10);
    std::cout << "Pi ~ " << global_sum << " with error "\
        << M_PI - global_sum << std::endl;
}

MPI_Finalize();

return 0;
}
```

# Collective communication rules

- **All** processes in the communicator must call the same collective function - trying to match an MPI\_Reduce with an MPI\_Recv is erroneous
- The calls to the collective functions must have compatible arguments - e.g. same destination process, datatype, count, reduction operation
- When a collective function returns the output to a single process, all other processes must still pass a (dummy) receive buffer argument - this could be just NULL
- Collective communications don't use the concept of a tag - they are matched **according to the order in which they are called**



# Order of collective calls matters

## Process 0

```
int a = 1, b = 2;  
int c, d;  
MPI_Reduce(&a, &c, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);  
MPI_Reduce(&b, &d, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
```

## Process 1

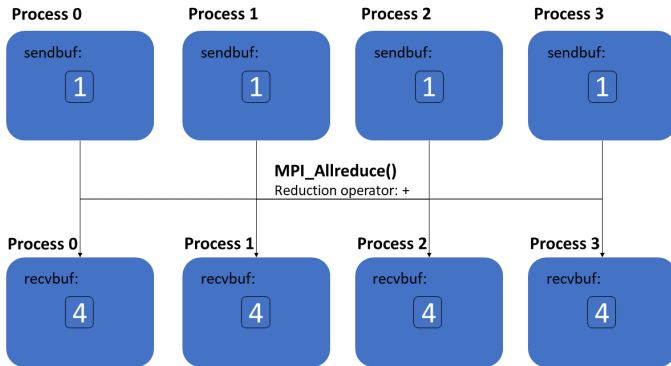
```
int a = 1, b = 2;  
int c, d;  
MPI_Reduce(&b, &d, ...);  
MPI_Reduce(&a, &c, ...);
```

## Process 2

```
int a = 1, b = 2;  
int c, d;  
MPI_Reduce(&a, &c, ...);  
MPI_Reduce(&b, &d, ...);
```

- We have two summation reductions:
  - `MPI_Reduce(&a, &c, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);`
  - `MPI_Reduce(&b, &d, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);`
- One would expect the reductions to sum variables a and b and store the results in variable c and d of rank 0
- **But** we would only get this outcome if **all processes called these functions in the same order**
- **Whereas now:**
  - The first collective reduction operator adds variables a of ranks 0 and 2 with variable b of rank 1 and stores in rank 0:  $c = 1 + 2 + 1$
  - The second collective reduction operator adds variables b of ranks 0 and 2 with variable a of rank 1 and stores in rank 0:  $d = 2 + 1 + 2$

# Reduction operation with result sent to all processes



## MPI\_Allreduce(): Reduction to all processes

```
int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count,\n                 MPI_Datatype datatype, MPI_Op op, MPI_Comm comm);
```

```
MPI_Allreduce(sendbuf, recvbuf, count, datatype, op, comm, ierror)\n  TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf\n  TYPE(*), DIMENSION(..) :: recvbuf\n  INTEGER, INTENT(IN) :: count\n  TYPE(MPI_Datatype), INTENT(IN) :: datatype\n  TYPE(MPI_Op), INTENT(IN) :: op\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Like the reduction operation, but now the result is returned to all processes in the group (hence the omission of a root process argument)

## MPI\_Wtime(): Taking timings

- In our previous example, we used the Unix shell command `time` to get timings for our program
- It is usually more helpful to time certain parts of our code, excluding time taken for input/output
- We are generally interested in **wall-clock time** (total elapsed time), rather than **CPU time**: the latter does not include idle time (e.g. a `MPI_Recv` waiting for a matching `MPI_Send`)
- MPI defines a convenient timer, `MPI_Wtime`:

```
double MPI_Wtime(void);
```

```
DOUBLE PRECISION MPI_Wtime()
```

- `MPI_Wtime` returns the time in seconds since some origin point in the past
- This origin is guaranteed to be constant within a process, but need not be the same for each process
- The function `MPI_Wtick()` returns the precision of `MPI_Wtime` in seconds (depends on the hardware)

# Example: synchronised timing

```
// Perform some computations

double start, end, elapsed, max, min, avg;

// Barrier makes sure all processes have finished earlier computation
MPI_Barrier(MPI_COMM_WORLD);
start = MPI_Wtime();

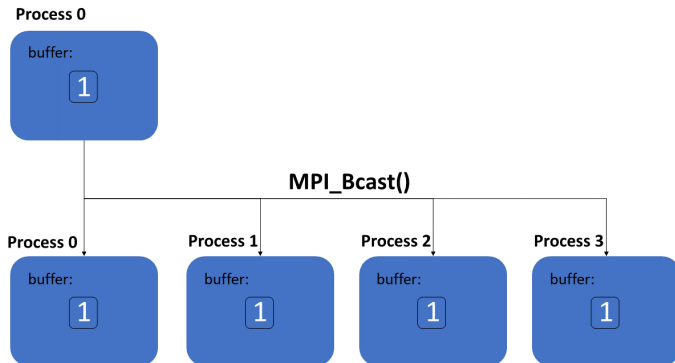
// Perform computation we want to time

end = MPI_Wtime();
elapsed = end - start;

// Reductions to get minimum, maximum and average timings
MPI_Reduce(&elapsed, &max, 1, MPI_DOUBLE, MPI_MAX, 0, MPI_COMM_WORLD);
MPI_Reduce(&elapsed, &min, 1, MPI_DOUBLE, MPI_MIN, 0, MPI_COMM_WORLD);
MPI_Reduce(&elapsed, &avg, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if (rank == 0)
{
    std::cout << "Maximum elapsed time: " << max << " s\nMinimum elapsed time: " << \
        min << " s\nAverage elapsed time: " << avg / size << " s" << std::endl;
}
```

# Broadcast operation



- **Broadcast** operation: Sends a message from a root process to all processes of the group including itself

## MPI\_Bcast(): Broadcast operation

```
int MPI_Bcast(void *buffer, int count, MPI_Datatype datatype, int root,\n              MPI_Comm comm);
```

```
MPI_Bcast(buffer, count, datatype, root, comm, ierror)\n  TYPE(*), DIMENSION(..) :: buffer\n  INTEGER, INTENT(IN) :: count, root\n  TYPE(MPI_Datatype), INTENT(IN) :: datatype\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- The root process sends the contents in buffer to all processes in the communicator comm
- By using a count greater than 1, we can send arrays of data

## Example: User input

- Most MPI implementations only allow process 0 in MPI\_COMM\_WORLD access to stdin
- This avoids having to decide which process gets which part of the input data

```
double x;
int n;
if (rank == 0)
{
    //Read user input
    std::cout << "Enter x: " << std::endl;
    std::cin >> x;
    std::cout << "Enter n: " << std::endl;
    std::cin >> n;
}

//Broadcast x and n to all processes
MPI_Bcast(&x, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

//Perform calculation using x and n
```



## Example: Input file reading

```
int file_size;
char *input_string

if (rank == 0)
{
    //Open file and get its size
    getFileAndSize(input_string, &file_size);
}

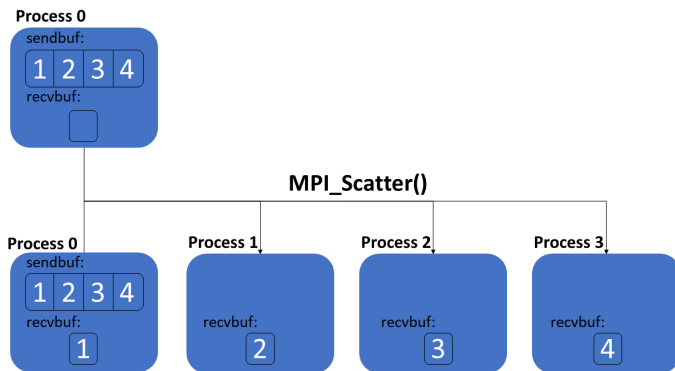
//Broadcast file size (count of chars)
MPI_Bcast(&file_size, 1, MPI_INT, 0, MPI_COMM_WORLD);

if (rank != 0)
{
    //Allocate enough memory to input buffers
    input_string = (char*) malloc(file_size * sizeof(char));
}

MPI_Bcast(input_string, input_size, MPI_CHAR, 0, MPI_COMM_WORLD);
```

- This technique should only be done for small input files
- MPI provides other ways for parallel input/output file operations (outside of this course's scope)

# Scatter operation



- **Scatter** operation: distributes array blocks among the processes in the communicator, in the order of their ranks
- In the above example, the array blocks have size 1


# MPI\_Scatter(): Scatter operation

```
int MPI_Scatter(const void *sendbuf, int sendcount, MPI_Datatype sendtype,\n               void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm);
```

```
MPI_Scatter(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, &\n            comm, ierror)\n    TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf\n    INTEGER, INTENT(IN) :: sendcount, recvcount, root\n    TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype\n    TYPE(*), DIMENSION(..) :: recvbuf\n    TYPE(MPI_Comm), INTENT(IN) :: comm\n    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- This sends the first block of sendcount elements in sendbuf of the root process to process 0, the next block of sendcount to process 1, etc. - i.e. we are scattering  $\text{sendcount} \times \text{nproc}$  elements in total
- The arguments sendbuf, sendcount and sendtype are only significant at the root
- On the root, the send and receive buffers should be different
- The receive count must match the corresponding send count on the root (unlike point-to-point communication)

# Gather operation



- **Gather** operation is the opposite of scatter
- It brings together blocks of data from all processors in the group and stacks it into a single array in the root process, in the order of the sending processes' ranks
- Be cautious of the memory requirements on the root, especially when sending large amounts of data from each process

## MPI\_Gather(): Gather operation

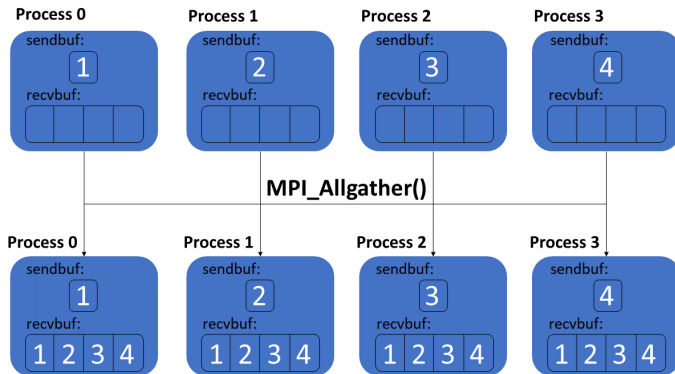
```
int MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype sendtype, \
              void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm);
```

```
MPI_Gather(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, &
           comm, ierror)
```

```
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf
INTEGER, INTENT(IN) :: sendcount, recvcount, root
TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype
TYPE(*), DIMENSION(..) :: recvbuf
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- The blocks of sendcount elements in sendbuf of all processes in comm are concatenated and stored in the recvbuf of the root process
- The arguments recvbuf, recvcount and recvtype are only significant at the root
- On the root, the send and receive buffers should be different
- The send count must match the corresponding receive count on the root (unlike point-to-point communication)

# Gather operation with result sent to all processes



## MPI\_Allgather(): Gather to all processes

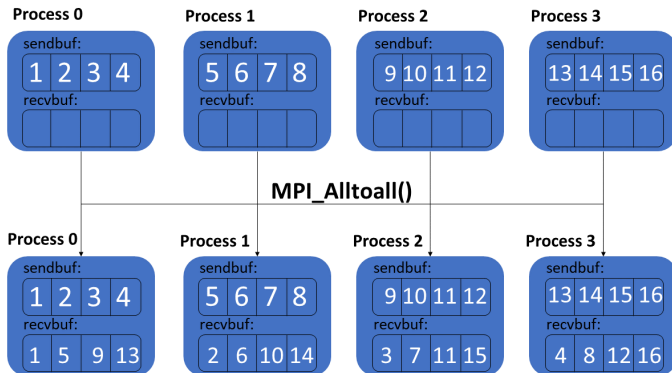
```
int MPI_Allgather(const void *sendbuf, int sendcount, MPI_Datatype,\n                 sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,\n                 MPI_Comm comm);
```

```
MPI_Allgather(sendbuf, sendcount, sendtype, recvbuf, recvcount,&\n              recvtype, comm, ierror)
```

```
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf\nINTEGER, INTENT(IN) :: sendcount, recvcount\nTYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype\nTYPE(*), DIMENSION(..) :: recvbuf\nTYPE(MPI_Comm), INTENT(IN) :: comm\nINTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Like the gather operation, but now the result is returned to all processes in the group (hence the omission of a root process argument)

# All-to-all operation



- The **All-to-All** operation is an extension to the all-gather operation, but now each process sends distinct data to each of the receivers
- Specifically, the  $j^{th}$  block of process  $i$  is placed as the  $i^{th}$  block in process  $j$  - like a matrix transpose operation



## MPI\_Alltoall(): All-to-all operation

```
int MPI_Alltoall(const void *sendbuf, int sendcount, MPI_Datatype,\n                sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype,\n                MPI_Comm comm);
```

```
MPI_Alltoall(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype,&\n            comm, ierror)\n    TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf\n    INTEGER, INTENT(IN) :: sendcount, recvcount\n    TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype\n    TYPE(*), DIMENSION(..) :: recvbuf\n    TYPE(MPI_Comm), INTENT(IN) :: comm\n    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- The send and receive buffers should be different to avoid aliasing
- The send and receive counts should be the same (unlike point-to-point communication)


## MPI\_IN\_PLACE: In-place communication

- Most of the collective communication functions introduced so far restrict the use of the same send and receive buffers to avoid aliasing
- We can still safely perform the operations in place, if we pass the special value `MPI_IN_PLACE` as one of the buffers
- For `MPI_Reduce` and `MPI_Gather`, `MPI_IN_PLACE` is used in the `sendbuf` argument at the root (or at all processes, if using `MPI_Allreduce`, `MPI_Allgather` or `MPI_Alltoall`)
- For `MPI_Scatter`, it is used in the `recvbuf` argument at the root
- For example, back to our code for estimating  $\pi$ , we could do the reduction in place, by replacing the `MPI_Reduce` call by:

```
if (rank == 0)
    MPI_Reduce(MPI_IN_PLACE, &sum, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
else
    MPI_Reduce(&sum, NULL, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

# Scatter/Gather count

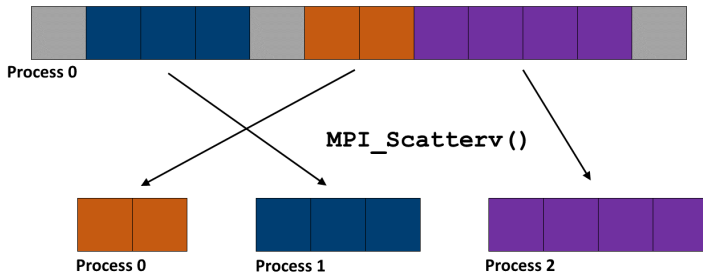
- So far, when scattering or gathering data, we used the same count of data elements in each process



- What if we want to scatter or gather an array with size that's not divisible by the number of processes?
- MPI also defines the variants `MPI_Scatterv`, `MPI_Gatherv`, `MPI_Allgatherv` (and `MPI_Alltoallv`), which allow having a different send/receive count on each rank
- These are even more flexible, in that they allow the scattered/gathered block data to be non-contiguous in memory

# Scatter operation with variable element count

```
sendcounts = [2, 3, 4]  
displs = [5, 1, 7]
```



# MPI\_Scatterv(): Variable count scatter operation

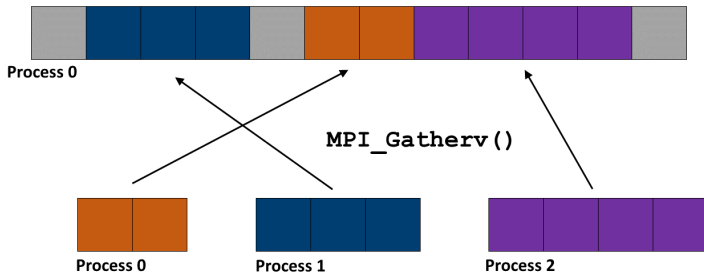
```
int MPI_Scatterv(const void *sendbuf, const int sendcounts[], const int displs[], MPI_Datatype sendtype,\n                void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm);
```

```
MPI_Scatterv(sendbuf, sendcounts, displs, sendtype, recvbuf, recvcount, recvtype, root, comm, ierror)\n  TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf\n  INTEGER, INTENT(IN) :: sendcounts(*), displs(*), recvcount, root\n  TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype\n  TYPE(*), DIMENSION(..) :: recvbuf\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- This extends the functionality of MPI\_Scatter by allowing irregular message sizes to be sent - hence the sendcount argument of MPI\_Scatter is now replaced by the array sendcounts
- The data can be sent in any order: the displacement (relative to sendbuf) of the block of data to be sent to process  $i$ , is the  $i^{th}$  entry of the array displs
- It is not allowed to read any data from the root more than once
- Fortran users: beware of conversion from array indices to ranks
- All other rules of MPI\_Scatter also apply here

# Gather operation with variable element count

```
recvcounts = [2, 3, 4]  
displs = [5, 1, 7]
```



## MPI\_Gatherv(): Variable count gather operation

```
int MPI_Gatherv(const void *sendbuf, int sendcount, MPI_Datatype sendtype,\n               void *recvbuf, const int recvcounts[], const int displs[],\n               MPI_Datatype recvtype, int root, MPI_Comm comm);
```

```
MPI_Gatherv(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, &\n            recvtype, root, comm, ierror)
```

```
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf
```

```
INTEGER, INTENT(IN) :: sendcount, recvcounts(*), displs(*), root
```

```
TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype
```

```
TYPE(*), DIMENSION(..) :: recvbuf
```

```
TYPE(MPI_Comm), INTENT(IN) :: comm
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Does exactly the opposite of MPI\_Scatterv
- Gathers a message of count `recvcounts[i]` from rank `i` and stores it into `recvbuf` of `root`, beginning at an offset of `displs[i]` elements
- The user needs to make sure that no location on the root is written more than once
- But the `recvbuf` need not be filled in a contiguous manner

## MPI\_Allgatherv(): Variable count gather to all processes

```
int MPI_Allgatherv(const void *sendbuf, int sendcount, MPI_Datatype\  
    sendtype, void *recvbuf, const int recvcounts[], const int\  
    displs[], MPI_Datatype recvtype, MPI_Comm comm);
```

```
MPI_Allgatherv(sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, &  
    recvtype, comm, ierror)
```

```
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf  
INTEGER, INTENT(IN) :: sendcount, recvcounts(*), displs(*)  
TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvtype  
TYPE(*), DIMENSION(..) :: recvbuf  
TYPE(MPI_Comm), INTENT(IN) :: comm  
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

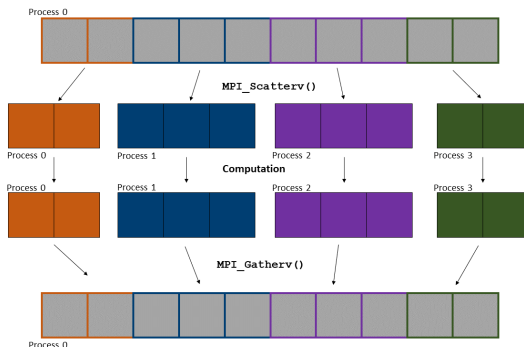
- Like MPI\_Gatherv, but now all processes receive the result, instead of just the root



# Example: Distributing and collecting data

Suppose

- We have a large array of data on rank 0 (potentially read from input file)
- We want to distribute this array (approximately) evenly to all processes
- Processes will perform some kind of computation on each element of the array
- And send their solutions back to rank 0 (e.g. to output results to file)



# Example: Distributing and collecting data

```
int arraylength = 100000;
double *global_array;
if (rank == 0)
{ // Set up data on rank 0 (e.g. by reading a file)
  getData(global_array);
}
// Compute the size of the array on every process
long istart = n * rank / size;
long iend = n * (rank + 1) / size;
int length = (int) (iend - istart);
// Get array sizes and displacements for communication
int lengths[size], displs[size];
MPI_Gather(&length, 1, MPI_INT, lengths, 1, MPI_INT, 0, MPI_COMM_WORLD);
if (rank == 0)
{
  displs[0] = 0;
  for (int i = 1; i < size; i++)
  {
    displs[i] = displs[i-1] + lengths[i-1];
  }
}
// Allocate enough memory to data buffers
double *array = (double*) malloc(length * sizeof(double));

// Distribute data from 0 to all processes
MPI_Scatterv(global_array, lengths, displs, MPI_DOUBLE, array, length, MPI_DOUBLE, 0, MPI_COMM_WORLD);


//Do some computation

// Collect updated array in rank 0
MPI_Gatherv(array, length, MPI_DOUBLE, global_array, lengths, displs, MPI_DOUBLE, 0, MPI_COMM_WORLD);
```


## More Point-To-Point

# Blocking and Non-blocking communication

- **Blocking communication:** send/receive operations which are guaranteed to only return after the corresponding buffers can be safely used



- **Non-blocking communication:** send/receive functions which start the corresponding operation but return immediately, without waiting for its completion (or freeing of buffers)



# Blocking vs. Non-blocking communication

## Blocking

- + Provides synchronisation: safe re-use of buffer is guaranteed with the return of the function call
- Can cause deadlocks
- The programmer should carefully order the send and receive calls - can be very difficult when using complex or dynamic communication patterns

## Non-blocking

- + Avoids deadlocks
- + Allows different communication operations to overlap (e.g. preventing serialisation)
- + Allows communication and computation operations to overlap
- Requires explicit checking for communication completion (through MPI-defined routines) before re-using send/receive buffers

- To provide more flexibility, MPI defines four communication modes for sending messages:
  - **Standard mode**: After the send call, the message can be either temporarily stored to a local buffer, or wait for a matching receive to be posted. It is up to the MPI implementation to decide, based on performance and memory considerations
  - **Buffered mode (B)**: If no matching receive is posted, the message is temporarily stored in some buffer space set by the programmer
  - **Synchronous mode (S)**: The send call waits for a matching receive to be posted and sends the message to the receive buffer
  - **Ready mode (R)**: The send call may be started only if the matching receive is already posted, otherwise the operation is erroneous

# Communication modes

- All sending communication modes are available in both blocking and non-blocking form (this course covers the blocking and standard non-blocking sends)

	Communication mode	MPI routine
<b>Blocking</b>	Standard	MPI_Send
	Buffered	MPI_Bsend
	Synchronous	MPI_Ssend
	Ready	MPI_Rsend
<b>Non-blocking</b>	Standard	MPI_Isend
	Buffered	MPI_Ibsend
	Synchronous	MPI_Issend
	Ready	MPI_Irsend

- There are two receive operations:
  - MPI\_Recv is the blocking receive
  - MPI\_Irecv is the non-blocking receive
- All sending routines can be matched with either of the two receives

## MPI\_Bsend(): Blocking buffered send

```
int MPI_Bsend(const void *buf, int count, MPI_Datatype datatype, int dest, \
              int tag, MPI_Comm comm);
```

```
MPI_Bsend(buf, count, datatype, dest, tag, comm, ierror)
  TYPE(*), DIMENSION(..), INTENT(IN) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Blocking send in buffered mode
- Can be used to avoid deadlocks
- But not ideal if memory is an issue
- Takes the same arguments as the standard send function MPI\_Send()
- The programmer should specify some buffer space on the sending process using MPI\_Buffer\_attach before calling MPI\_Bsend



## MPI\_Buffer\_attach(): Setting a buffer

```
int MPI_Buffer_attach(void *buffer, int size);
```

```
MPI_Buffer_attach(buffer, size, ierror)  
  TYPE(*), DIMENSION(..), ASYNCHRONOUS :: buffer  
  INTEGER, INTENT(IN) :: size  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Provides a buffer, of size size (in bytes) in the user's memory, to be used for buffer-mode outgoing messages
- If not specified, a zero-sized buffer is associated with the process
- Only one buffer can be attached to a process at a time
- So the size of the buffer should be big enough to cover all pending buffered sends from the process

## MPI\_Buffer\_detach(): Removing a buffer

```
int MPI_Buffer_detach(void *buffer_addr, int *size)
```

```
MPI_Buffer_detach(buffer_addr, size, ierror)  
  USE, INTRINSIC :: ISO_C_BINDING, ONLY : C_PTR  
  TYPE(C_PTR), INTENT(OUT) :: buffer_addr  
  INTEGER, INTENT(OUT) :: size  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Detaches the buffer currently associated with this process
- Returns the buffer address and size
- Blocks until all messages in the buffer have been transmitted
- Once this call returns, the programmer can safely reuse or deallocate the space taken by the buffer

- The call to `MPI_Pack_size(count, datatype, comm, &data_size)` calculates an upper bound for the amount of space (in bytes) needed to pack a message, and returns this value in `data_size`
- The constant `MPI_BSEND_OVERHEAD` gives an upper bound on the space needed for additional information (e.g. tag, destination, communicator)

## Example: buffered sends

- So, if we wanted to use a buffer-mode send to transmit an integer from rank 0 to all other processes, we could do the following:

```
int comm_size, rank;
MPI_Comm_size(MPI_COMM_WORLD, &comm_size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);

if (rank == 0)
{
    int a = 10; // integer to be sent

    // Calculate buffer size in bytes
    int data_size;
    MPI_Pack_size(1, MPI_INT, MPI_COMM_WORLD, &data_size);
    int buf_size = (comm_size - 1) * (data_size + MPI_BSEND_OVERHEAD);

    int* buf = (int*) malloc(buf_size);
    MPI_Buffer_attach(buf, buf_size);

    for (int i = 1; i < comm_size; i++)
    {
        MPI_Bsend(&a, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
    }

    //Do other work

    //When all sends are expected to have completed, detach and free the buffer
    MPI_Buffer_detach(&buf, &buf_size);
    free(buf);
}
```

## MPI\_Ssend(): Blocking synchronous send

```
int MPI_Ssend(const void *buf, int count, MPI_Datatype datatype, int dest, \
              int tag, MPI_Comm comm);
```

```
MPI_Ssend(buf, count, datatype, dest, tag, comm, ierror)
  TYPE(*), DIMENSION(..), INTENT(IN) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Blocking send in synchronous mode
- Takes the same arguments as MPI\_Send()
- Can be used in place of standard sends to check whether a program is *safe*: i.e. if it does not depend on message buffering for its completion
- Safe programs are more portable because they don't depend on the communication protocol used or the amount of buffer space available

## MPI\_Rsend(): Blocking ready send

```
int MPI_Rsend(const void *buf, int count, MPI_Datatype datatype, int dest, \
              int tag, MPI_Comm comm);
```

```
MPI_Rsend(buf, count, datatype, dest, tag, comm, ierror)
  TYPE(*), DIMENSION(..), INTENT(IN) :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Blocking send in ready mode
- Requires careful programming! It should only be used when it's certain that the receive has already been posted, otherwise the operation is erroneous and the outcome is undefined
- A ready send could improve performance by saving the overhead of checking for a matching receive call
- Takes the same arguments as MPI\_Send()

## MPI\_Isend(): Non-blocking send

```
int MPI_Isend(const void *buf, int count, MPI_Datatype datatype, int dest, \
             int tag, MPI_Comm comm, MPI_Request *request);
```

```
MPI_Isend(buf, count, datatype, dest, tag, comm, request, ierror)
  TYPE(*), DIMENSION(..), INTENT(IN), ASYNCHRONOUS :: buf
  INTEGER, INTENT(IN) :: count, dest, tag
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  TYPE(MPI_Request), INTENT(OUT) :: request
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Initiates a non-blocking send in standard mode
- Arguments are as in the blocking sends, except from the final request argument
- MPI\_Request is used as an identifier for the immediate send
- It can then be used in completion calls, to request information about the send operation status (before re-using the send buffer)

## MPI\_Irecv(): Non-blocking receive

```
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source,\n              int tag, MPI_Comm comm, MPI_Request *request);
```

```
MPI_Irecv(buf, count, datatype, source, tag, comm, request, ierror)\n  TYPE(*), DIMENSION(..), ASYNCHRONOUS :: buf\n  INTEGER, INTENT(IN) :: count, source, tag\n  TYPE(MPI_Datatype), INTENT(IN) :: datatype\n  TYPE(MPI_Comm), INTENT(IN) :: comm\n  TYPE(MPI_Request), INTENT(OUT) :: request\n  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Initiates a non-blocking receive
- Arguments are as in the blocking receive, except from the status argument which is now replaced by the request argument
- This can then be used in completion calls, to request information about the receive operation status (before using the receive buffer)



# Communication completion operations

- With non-blocking communication, it is essential to ensure that communication has completed before re-using the send/receive buffers
- There are two types of completion testing functions:
  - **Wait** returns when the operation is complete - this is blocking: it waits until the operation is complete
  - **Test** returns a flag which is set to true if the operation has completed, or false otherwise - this is non-blocking: it returns immediately
- Both of these take the request argument passed in the non-blocking send/receive call to identify which operation to wait on or test

## MPI\_Wait(): Wait for completion

```
int MPI_Wait(MPI_Request *request, MPI_Status *status);
```

```
MPI_Wait(request, status, ierror)  
  TYPE(MPI_Request), INTENT(INOUT) :: request  
  TYPE(MPI_Status) :: status  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Waits until the operation identified by request is complete
- Recall: a send/receive operation is complete when the send/receive buffers are safe to be re-used
- When the call returns, it also deallocates the request and sets request to MPI\_REQUEST\_NULL (a null request handle)
- If used for a receive operation, the call also returns information on the completed operation in status (see slide 42)

## MPI\_Test(): Test for completion

```
int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
```

```
MPI_Test(request, flag, status, ierror)  
  TYPE(MPI_Request), INTENT(INOUT) :: request  
  LOGICAL, INTENT(OUT) :: flag  
  TYPE(MPI_Status) :: status  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- Tests to find out if the operation identified by request is complete
- If so, it returns flag = true, it deallocates the request and sets request to MPI\_REQUEST\_NULL
- Otherwise it returns flag = false
- If used for a receive operation, the call also returns information on the completed operation in status (see slide 42)

## Example: Non-blocking point-to-point communication

```
if (rank == 0)
{
    MPI_Isend(a, 10, MPI_DOUBLE, 1, tag, comm, request);

    // Perform some other computation without altering variable a

    MPI_Wait(request, status);
    // Can now safely change variable a
}
else (if rank == 1)
{
    MPI_Irecv(a, 15, MPI_DOUBLE, 0, tag, comm, request);

    // Perform some other computation without using variable a

    MPI_Wait(request, status);
    // Can now safely read and use a
}
```

# Multiple completions

- It is possible to test or wait for the completion of a number of non-blocking operations at the same time
- MPI defines variants of the MPI\_Wait and MPI\_Test for waiting/testing whether all/some/any of the operations in a given list have completed
- MPI\_Waitall/MPI\_Testall: Waits/tests for completion of **all** operations
- MPI\_Waitsome/MPI\_Testsome: Waits/tests for completion of **at least one** operation
- MPI\_Waitany/MPI\_Testany: Waits/tests for completion of **any one** of the listed operations

- It is good practice to initialise requests to `MPI_REQUEST_NULL`:
  - If a process accidentally calls a communication completion function with an uninitialised request handle, the outcome is undefined (possible a segfault)
  - But communication functions called with `MPI_REQUEST_NULL` are simply null operations
- If we want to test for completion without deallocating the request, we can use `MPI_Request_get_status(request, flag, status)`
- We can deallocate requests and set them to `MPI_REQUEST_NULL` by calling `MPI_Request_free(request)`
- **Caution:** This should only be used if we know communication has completed, for example if we called `MPI_Request_get_status` or if we match `MPI_Isend` with `MPI_Recv`. If an error occurs during a communication after the request object has been freed, the error code cannot be returned to the user

- Probing: MPI allows us to check for incoming messages, without actually receiving them
- This also provides information that is useful for knowing *how* to receive them (e.g. message type, count etc.)
- There is a blocking and non-blocking probing function, MPI\_Probe and MPI\_Iprobe

## MPI\_Probe(): Blocking probe

```
int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status);
```

```
MPI_Probe(source, tag, comm, status, ierror)
```

```
INTEGER, INTENT(IN) :: source, tag
```

```
TYPE(MPI_Comm), INTENT(IN) :: comm
```

```
TYPE(MPI_Status) :: status
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- MPI\_Probe waits until an incoming message with matching source, tag and communicator comm has been found and returns its status
- The returned status argument is the same value that would have been returned by MPI\_Recv



## MPI\_Iprobe(): Non-blocking (immediate) probe

```
int MPI_Iprobe(int source, int tag, MPI_Comm comm, int *flag, MPI_Status *status)
```

```
MPI_Iprobe(source, tag, comm, flag, status, ierror)
```

```
INTEGER, INTENT(IN) :: source, tag
```

```
TYPE(MPI_Comm), INTENT(IN) :: comm
```

```
LOGICAL, INTENT(OUT) :: flag
```

```
TYPE(MPI_Status) :: status
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- MPI\_Iprobe is like MPI\_Probe, but it is non-blocking
- It checks to see if there is an incoming message with matching source, tag and communicator comm and returns immediately
- If such a message can be found, it returns flag = true and fills the status argument, otherwise it returns flag = false and leaves status undefined

## Error handling

- So far, we've completely ignored the error code returned from MPI routines
- This is the return value of the MPI function in C/C++, or the `ierror` argument in Fortran
- The reason for this is that by default, MPI sets the error handler to `MPI_ERRORS_ARE_FATAL`, which means that when an error occurs, the program aborts
- If we want to know more about the error and let the program continue, we can reset the error handler to `MPI_ERRORS_RETURN` by using `MPI_Comm_set_errhandler(MPI_COMM_WORLD, MPI_ERRORS_RETURN);`
- This should be used with caution: if unstopped, errors could propagate and cause more errors further down

# Error codes

- MPI defines MPI\_SUCCESS, which is returned when the routine finishes with no error
- All other return codes are implementation-defined
- But MPI defines a small set of error classes to divide the error codes into categories depending on the type of the error, including:

Error class	Meaning
MPI_SUCCESS	No error
MPI_ERR_BUFFER	Invalid buffer pointer
MPI_ERR_COUNT	Invalid count argument
MPI_ERR_TYPE	Invalid datatype argument
MPI_ERR_TAG	Invalid tag argument
MPI_ERR_COMM	Invalid communicator
MPI_ERR_RANK	Invalid rank
MPI_ERR_REQUEST	Invalid request handle
MPI_ERR_ROOT	Invalid root
MPI_ERR_OP	Invalid operation

# Retrieving errors

- The MPI routine `MPI_Error_class` converts each error code into the corresponding error class

```
int MPI_Error_class(int errorcode, int *errorclass);
```

```
MPI_Error_class(errorcode, errorclass, ierror)  
  INTEGER, INTENT(IN) :: errorcode  
  INTEGER, INTENT(OUT) :: errorclass  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- To obtain a human readable description of the error, we can use the function `MPI_Error_string`, where the output argument `resultlen` gives the length of the error string

```
int MPI_Error_string(int errorcode, char *string, int *resultlen);
```

```
MPI_Error_string(errorcode, string, resultlen, ierror)  
  INTEGER, INTENT(IN) :: errorcode  
  CHARACTER(LEN=MPI_MAX_ERROR_STRING), INTENT(OUT) :: string  
  INTEGER, INTENT(OUT) :: resultlen  
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

## Example: Error handling

```
// Change default behaviour of error handler, so that error codes are returned
MPI_Comm_set_errhandler(MPI_COMM_WORLD, MPI_ERRORS_RETURN);

// Error: send to an invalid rank of index -1
int errcode;
errcode = MPI_Send(&x, 1, MPI_DOUBLE, -1, 0, MPI_COMM_WORLD);
if (errcode != MPI_SUCCESS)
{
    /* The error code will either be MPI_SUCCESS, or some error code that falls
    under an error class, which we can retrieve */
    int errclass;
    MPI_Error_class(errcode, &errclass);
    if (errclass == MPI_ERR_RANK)
    {
        // The string must be at least MPI_MAX_ERROR_STRING characters long
        char errstring[MPI_MAX_ERROR_STRING];
        int resultlen;
        MPI_Error_string(errcode, errstring, &resultlen);
        std::cout << errstring << std::endl;

        // Then either find some fix and continue the application or call MPI_Abort
    }
}
```

## Common mistakes

- **Mistake:** Doing things before `MPI_Init` or after `MPI_Finalize`  
**Cause:** The MPI standard does not say much about the region before `MPI_Init` or after `MPI_Finalize` - not even how many processes are running. It does require that rank 0 of `MPI_COMM_WORLD` returns from `MPI_Finalize`, but that's only if it hasn't been terminated before (e.g. aborted). Any other behaviour will be implementation dependent.  
**Solution:** If possible, avoid doing any tasks (other than specific MPI routines like `MPI_Get_version`) before initialising or after finalising MPI
- **Mistake:** Using a different MPI implementation in compile, link or run time  
**Cause:** Different implementations make different choices in the way they are written  
**Effect:** Link-time errors, MPI program failing to start, or multiple serial versions running instead of multiple communicating ones



## Deadlock

When processes are waiting for an event that will never happen and the program hangs

- **Mistake:** Posted send and receive operations do not match  
**Cause:** This might be due to the operations having different tags, or datatypes, or the receive buffer being shorter than the message size, for example  
**Effect:** Deadlock  
**Solution:** Make sure all send-receive pairs follow the rules of slide 41
- **Mistake:** Calls to send and receive posted in wrong order  
**Cause:** For example, posting the receives first, on both processes  
**Effect:** Deadlock  
**Solution:** Use one of: Non-blocking operations, blocking buffered send, (MPI\_Bsend), or MPI\_Sendrecv

# Mistakes using point-to-point communication

- **Mistake:** Not making programs safe

**Cause:** We are relying on message buffering when sending with `MPI_Send`

**Effect:** Program works for small message sizes, but will cause deadlock for larger messages

**Solution:** Use one of: Non-blocking operations, blocking buffered send, (`MPI_Bsend`), or `MPI_Sendrecv`

## Race condition

When two or more operations attempt to read/write the same location in memory at the same time, resulting in corrupted data

- **Mistake:** Re-using send/receive buffers of non-blocking operations before the send/receive is complete

**Cause:** Non-blocking operations return immediately, without waiting for communication completion

**Effect:** Results change if the program is run multiple times

**Solution:** Use a communication completion testing routine, like `MPI_Wait` or `MPI_Test`

# Mistakes using point-to-point communication

- **Mistake:** Assuming that the dealing of messages is fair

**Cause:** Messages are non-overtaking, but non necessarily fair. If two messages are sent **from the same source to the same destination** and a receive operation matches both messages, the message posted first will be received first (the same applies for 2 receives matching the same message: the receive posted first will receive the message). **But** messages might be overtaken by matching messages coming from another process, even if they were sent later

**Effect:** Wrong results if sends and receives are unintentionally matched, potentially different every time we run the program

**Solution:** Avoid using `MPI_ANY_SOURCE` or `MPI_ANY_TAG` too much. If using non-blocking communication, one can use `MPI_Testsome` and handle all complete communications

# Mistakes using collective communication

- **Mistake:** Not calling collective operations on all processes in the communicator, or calling them in different orders  
**Effect:** Deadlock, or wrong results  
**Solution:** Take care, especially when using if statements
- **Mistake:** Expecting all processes to return from collective calls at (roughly) the same time  
**Solution:** If it is essential that processes synchronise, use `MPI_Barrier`
- **Mistake:** Trying to match an `MPI_Bcast` with an `MPI_Recv`  
**Cause:** Even though acts as a multiple-send, `MPI_Bcast` is a collective operation  
**Solution:** Call `MPI_Bcast` on all processes in the communicator, with no call to `MPI_Recv`
- **Mistake:** Using the same input and output buffer (e.g. in `MPI_Reduce`, `MPI_Gather` etc)  
**Cause:** This is not allowed by the MPI standard, to avoid aliasing  
**Effect:** This may work in some cases, but may fail in others  
**Solution:** Use `MPI_IN_PLACE` (see slide 90)

## Performance Tips

- Use collective, rather than point-to-point communication, where possible
- Instead of sending lots of small messages, it will be faster to send less messages of a larger size (e.g. using `count`  $\gg 1$ , or MPI-derived datatypes)
- Don't use more processes than cores
- If the application is not big or complex enough, consider running in serial first, with compiler optimisations turned on

## Further topics

- MPI is a large and evolving standard
- We have only covered a small part of it
- There are hundreds of MPI routines which provide additional functionality, suited to different applications and needs



# User-defined data types

- We have introduced and used basic MPI data types (slide 43)
- It is possible to combine data into a single, user-defined datatype
- The advantage is that we can send this data using just one message
- There are various ways to do this - we will very briefly mention one of them

# User-defined data types in practice

- If we want to combine an array of 10 doubles, stored in variable x and an integer stored in y into a new data type
- And then broadcast from 0 to all other processes

```
MPI_Datatype newtype;

int array_of_blocklengths[2] = {10, 1};
MPI_Datatype array_of_types[2] = {MPI_DOUBLE, MPI_INT};
MPI_Aint x_addr, y_addr;
MPI_Aint array_of_displacements[2]; // Displacement of elements from first element (in bytes)

array_of_displacements[0] = 0;
MPI_Get_address(&x, &x_addr); // Returns address of x on the system
MPI_Get_address(&y, &y_addr);
array_of_displacements[1] = y_addr - x_addr;

// Build the derived data type:
MPI_Type_create_struct(2, array_of_blocklengths, array_of_displacements, array_of_types, &newtype);
// Before using the datatype, we should first commit it:
MPI_Type_commit(&newtype);

MPI_Bcast(&x, 1, newtype, 0, MPI_COMM_WORLD);

// After using it, free any additional storage used to create the data type
MPI_Type_free(&newtype);
```

# User-defined communicators

- **Recall:** A communicator is a collection of processes that can send messages to each other
- `MPI_COMM_WORLD` is the default communicator consisting of all processes in the program
- `MPI_COMM_SELF` is also provided - this only includes the process itself (not as useful)
- Within their `MPI_COMM_SELF`, all processes have rank 0
- The user can also define their own communicators
- This is particularly useful when the program has multiple levels of parallelism, i.e. can be split into a number of further parallelisable sub-tasks
- Using user-defined communicators insures that data relating to one sub-task are not accidentally sent to another

# User-defined communicators

- New communicators can only be built upon existing communicators
- There are different ways to construct new communicators
- `MPI_Comm_split` is a collective call which partitions an existing communicator into disjoint sub-communicators

```
int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm *newcomm);
```

```
MPI_Comm_split(comm, color, key, newcomm, ierror)
```

```
TYPE(MPI_Comm), INTENT(IN) :: comm
```

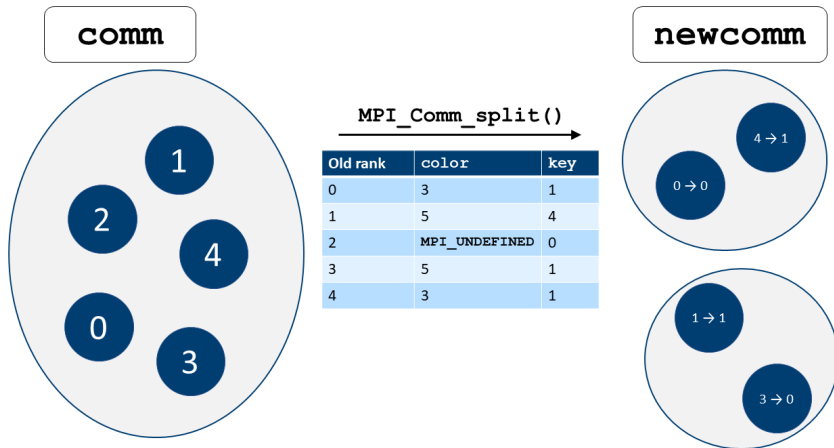
```
INTEGER, INTENT(IN) :: color, key
```

```
TYPE(MPI_Comm), INTENT(OUT) :: newcomm
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- `comm` is the existing communicator
  - `color` is a non-negative integer which determines to which new communicator each process will belong (or `MPI_UNDEFINED` if process is joining none)
  - `key` determines the ordering (ranks) within the new communicators - i.e. process with smallest key receives rank 0 etc.
  - `newcomm` is the new communicator
- We can remove this by calling `MPI_Comm_free(newcomm)`

# Splitting a communicator



- The MPI standard does not require that all nodes can perform I/O (even though our examples have assumed that all processes can do output and process 0 can also do input)
- To find which processes can do I/O, one can call `MPI_Comm_get_attr(MPI_COMM_SELF, MPI_IO, attribute_val, flag)`
- `attribute_val` is a pointer, which will hold the returned value:

Returned value	Meaning
<code>MPI_ANY_SOURCE</code>	all ranks can do I/O
Processes' rank	the current process can do I/O (but not all)
Another rank	some rank can do I/O
<code>MPI_PROC_NULL</code>	no rank can do I/O

- So we can check our assumption that rank 0 can do I/O by running this function on rank 0 and checking for a return value of either `MPI_ANY_SOURCE` or 0
- MPI also provides functions for parallel I/O (e.g. to read from and write to a file in parallel), which are outside of the scope of this course

# Using MPI with OpenMP

- Need even more speed-up?
- Nothing stops us from using thread-based parallelisation (e.g. OpenMP) within a node and MPI between nodes, in a *hybrid* approach
- **But** we should do this in a *thread safe* manner
- The call to `MPI_Init` should be replaced by `MPI_Init_thread`, which initialises the thread environment to the required level of thread support, if possible (depends on the MPI implementation)

```
int MPI_Init_thread(int *argc, char ***argv, int required, int *provided);
```

```
MPI_Init_thread(required, provided, ierror)
```

```
    INTEGER, INTENT(IN) :: required
```

```
    INTEGER, INTENT(OUT) :: provided
```

```
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

Thread support level	Meaning
<code>MPI_THREAD_SINGLE</code>	Only one thread
<code>MPI_THREAD_FUNNELED</code>	Multiple threads - a master thread makes all MPI calls
<code>MPI_THREAD_SERIALIZED</code>	Multiple threads - only call MPI one at a time
<code>MPI_THREAD_MULTIPLE</code>	Multiple threads - may all call MPI with no restrictions

- `MPI_Init_thread` should be called once on every process
- The thread which calls it becomes the main thread of that process
- The two routines below are callable from any thread at any time (unlike most MPI routines):
  - `MPI_Is_thread_main(int *flag)` determines whether the current thread is the main thread (the one that called `MPI_Init_thread`)
  - `MPI_Query_thread(int *provided)` returns the current level of thread support - i.e. the same as the one provided by `MPI_Init_thread`



In this course, we

- gave a brief introduction to parallel programming
- introduced the message passing interface and its syntax
- looked at point-to-point communication for sending messages from one process to another (and the potential issues arising from mismatched/hanging calls)
- looked at MPI functions for collective communication between all processes in a communicator, which simplified the programming, increased performance and eliminates some of the risks associated with point-to-point communication
- introduced other modes of sending messages (e.g. blocking vs non-blocking), which allow the overlapping of calculation and communication
- looked at the errors MPI functions may return,
- as well as mistakes which are not caught by MPI
- mentioned some further topics (user-defined data types, user-defined communicators and how to use MPI with OpenMP)

- There is nothing wrong with using this small subset of MPI
- Keeping it simple may prove to be less prone to bugs and equally fast as something more sophisticated
- But there might be applications which require further functionality
- You can always refer to the MPI standard, found at <https://www.mpi-forum.org/docs/>, or the books suggested in slide 3
- Happy parallelising! :)